

Q1: TensorFlow vs. PyTorch

Aspect	TensorFlow	PyTorch
Execution Model	Uses static computation graphs by default (<code>@tf.function</code>)—good for optimization and deployment	Eager execution by default—operations are executed immediately, more intuitive and “Pythonic”
Ease of Debugging	Harder with static graphs, though <code>tf.debugging</code> tools exist	Easier with dynamic graphs—you can use standard Python debugging tools like <code>pdb</code>
Model Deployment	Robust ecosystem for deployment: TensorFlow Lite (mobile), TensorFlow Serving (web), TensorFlow.js	Deployment options improving (e.g., TorchServe), but not as mature or widespread
Visualization	TensorBoard: powerful suite for visualizing training, hyperparameters, histograms, etc.	Basic visualization tools; community tools like Weights & Biases or TensorBoard with <code>torch.utils</code>
Syntax/Code Style	Verbose, especially in older versions; higher learning curve for beginners	More readable and concise—more like writing NumPy code
Community & Adoption	Backed by Google; widely used in industry production systems	Strong academic/research adoption; backed by Meta AI (formerly Facebook AI)

When to choose:

- Choose **PyTorch** if you're in **research, rapid prototyping**, or prefer readable, native-Python code.
- Choose **TensorFlow** if you're building **large-scale production pipelines**, want **tight integration with Google Cloud**, or need mobile/web deployment via TF Lite or TF.js..

Q2: Jupyter Notebook Use Cases in AI

1. **Exploratory Data Analysis (EDA)** In AI, understanding your dataset is critical. Jupyter allows:
 - Writing Python code alongside output plots (e.g., using Matplotlib, Seaborn, or Plotly)
 - Using Markdown cells to describe findings

- Running statistical summaries (mean, variance, null counts) interactively ✨
Example: You might build histograms of feature distributions, visualize missing values, or generate correlation heatmaps—all in a single notebook.

2. Interactive Model Development

- You can build, tweak, and test machine learning models (e.g., with scikit-learn, Keras, PyTorch) step by step.
- Easy to track losses, visualize confusion matrices, and test different hyperparameters in-line.

Q3: How spaCy Enhances NLP vs. Basic Python String Operations

Basic string operations like `.split()`, `.replace()`, `.lower()` are good for surface-level manipulation but lack linguistic understanding. Here's how **spaCy takes it to the next level**:

Feature	Basic Python	spaCy
Tokenization	<code>.split()</code> splits on whitespace	Recognizes punctuation, contractions, and special tokens (e.g., “can’t” → “ca”, “n’t”)
Part-of-Speech Tagging	Not available	<code>token.pos_</code> gives grammatical role (noun, verb, etc.)
Named Entity Recognition	Not available	Detects named entities (e.g., PERSON, ORG, DATE) with <code>doc.ents</code>
Lemmatization	Manual or NLTK	Built-in <code>.lemma_</code> gives base forms (e.g., “running” → “run”)
Dependency Parsing	Not supported	Parses syntactic structure: subject, object, modifiers
Word Vectors / Similarity	Not available	Pre-trained vectors (via <code>en_core_web_md</code>) for semantic similarity calculations

Target Applications

Aspect	Scikit-learn	TensorFlow
Focus	Classical ML algorithms (SVMs, Random Forests, Logistic Regression, etc.)	Deep learning and neural networks (CNNs, RNNs, Transformers)
Use Case Fit	Great for structured/tabular data, quick experimentation, smaller datasets	Ideal for image/audio processing, natural language tasks, large datasets, and production
Model Complexity	Pre-built, shallow learning models, limited support for custom architectures	Highly customizable models—from simple MLPs to cutting-edge transformer stacks

Ease of Use for Beginners

Aspect	Scikit-learn	TensorFlow
Learning Curve	Beginner-friendly—clear API, no need for GPUs or complex tensor operations	Steeper curve, especially for low-level API (though Keras helps simplify that)
API Design	Consistent and elegant— <code>.fit()</code> , <code>.predict()</code> , <code>.score()</code> for almost everything	More layered—can be simple with Keras or deep with core TensorFlow APIs
Environment Setup	Lightweight; doesn't require CUDA/GPU installations	Heavier setup if doing GPU-based training

Community Support

Aspect	Scikit-learn	TensorFlow
Maturity	Established since 2010; extremely stable	Launched in 2015; very active development
Documentation	Highly readable, beginner-friendly, lots of example datasets	Extensive docs and tutorials, especially for deep learning workflows
Community Contributions	Huge ecosystem; often used alongside Pandas, NumPy, and Matplotlib	Thriving GitHub, Stack Overflow, and TensorFlow Hub community

Enterprise Adoption

Common in traditional analytics or smaller-scale ML pipelines

Widely used in production systems, especially with cloud platforms and edge devices

1. Ethical Considerations

Potential Biases in MNIST Model:

- **Dataset bias:** MNIST contains primarily Western-style handwritten digits
- **Cultural bias:** Writing styles vary across cultures and regions
- **Demographic bias:** Limited representation of different age groups, handwriting abilities

Potential Biases in Amazon Reviews Model:

- **Language bias:** English-only analysis excludes non-English speakers
- **Platform bias:** Amazon reviews may not represent all consumer demographics
- **Temporal bias:** Sentiment patterns may change over time
- **Product category bias:** Different products may have different review patterns